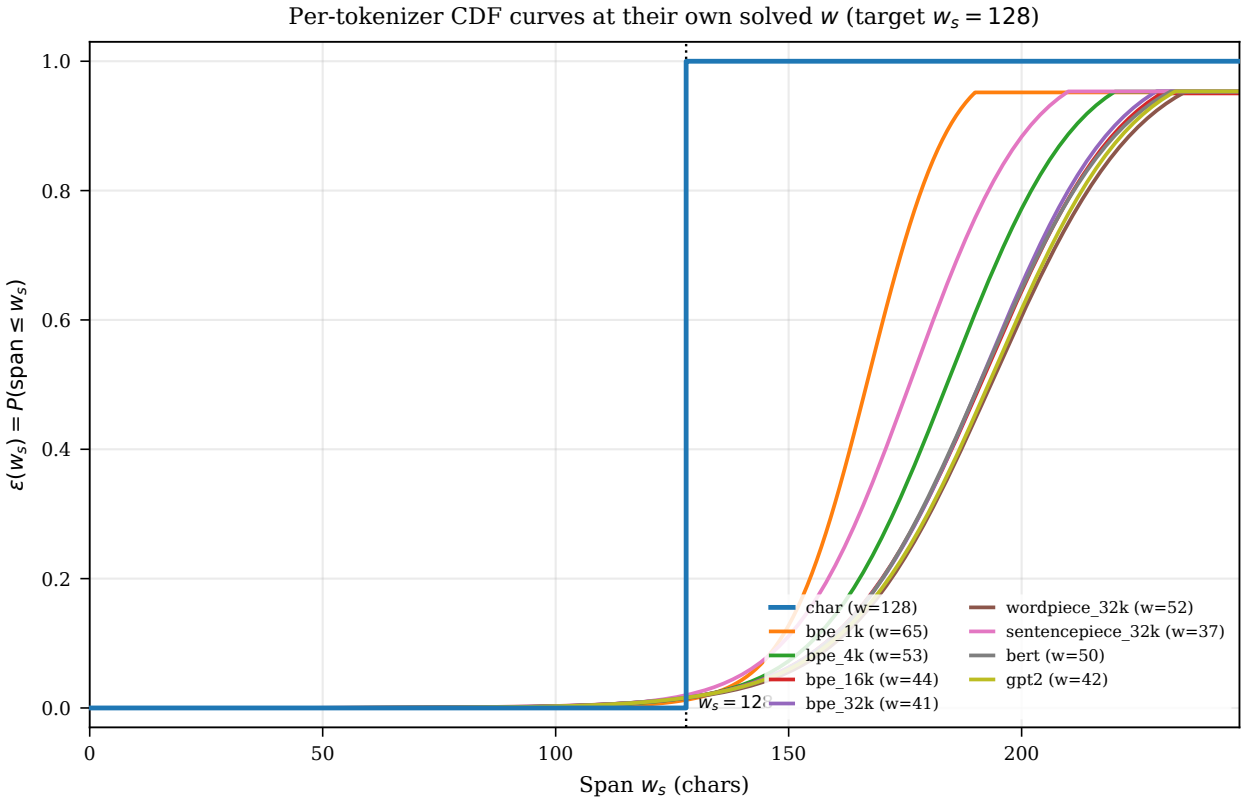
Per-tokenizer CDF curves at their own solved w (target $w_s = 128$)
$\varepsilon(w_s) = P(\text{span} \le w_s)$
Span w_s (chars)
1.0
0.8
0.6
0.4
0.2
0.0
0
50
100
150
200
$w_s = 128$
char (w=128)
bpe_1k (w=65)
bpe_4k (w=53)
bpe_16k (w=44)
bpe_32k (w=41)
wordpiece_32k (w=52)
sentencepiece_32k (w=37)
bert (w=50)
gpt2 (w=42)