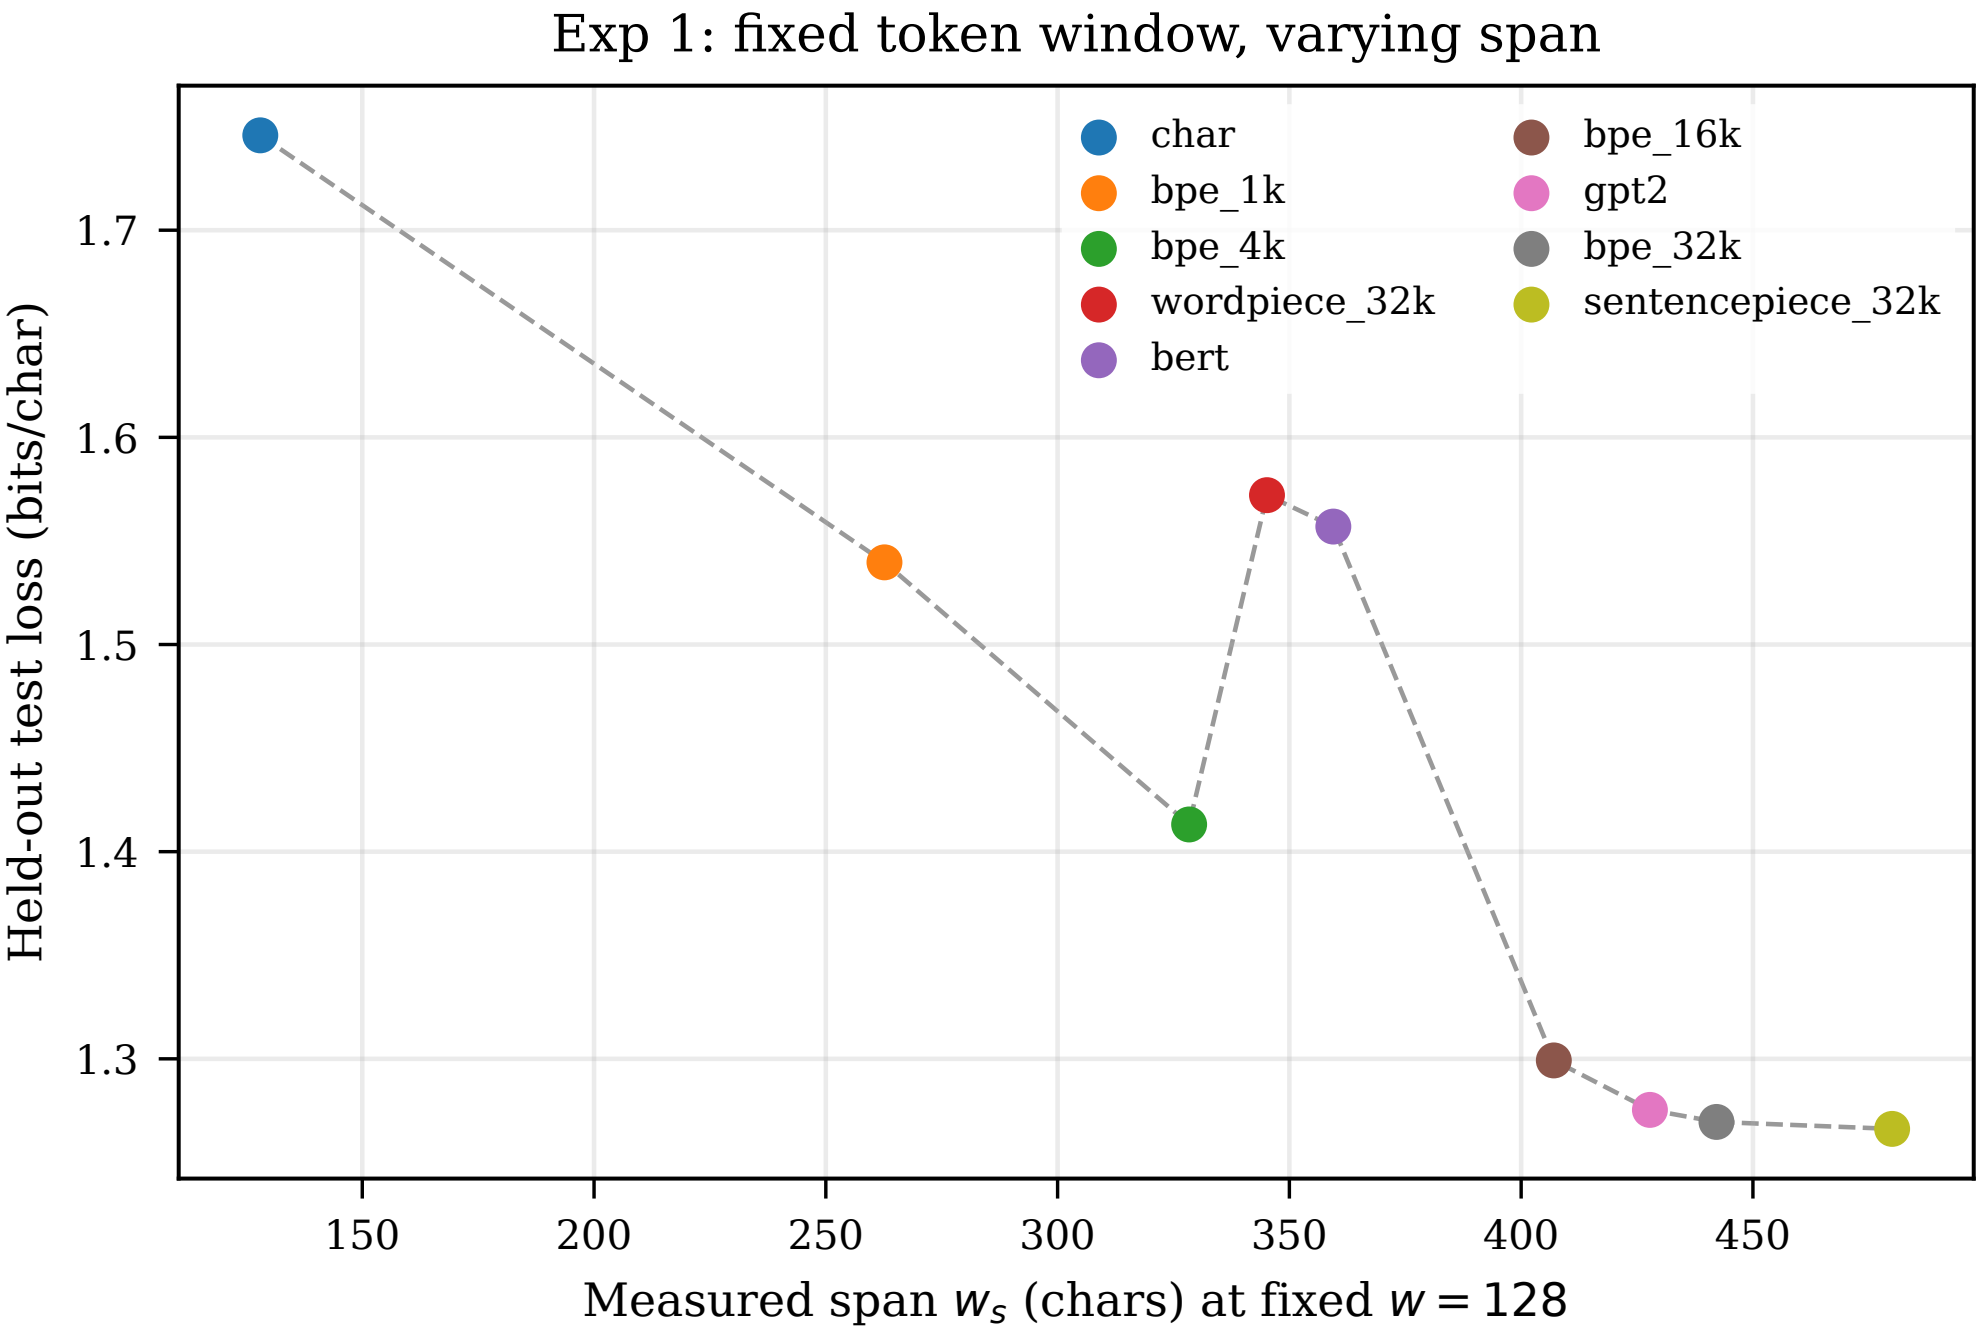
Exp 1: fixed token window, varying span
char
bpe_1k
bpe_4k
wordpiece_32k
bert
bpe_16k
gpt2
bpe_32k
sentencepiece_32k
Held-out test loss (bits/char)
1.7
1.6
1.5
1.4
1.3
150
200
250
300
350
400
450
Measured span w_s (chars) at fixed $w = 128$