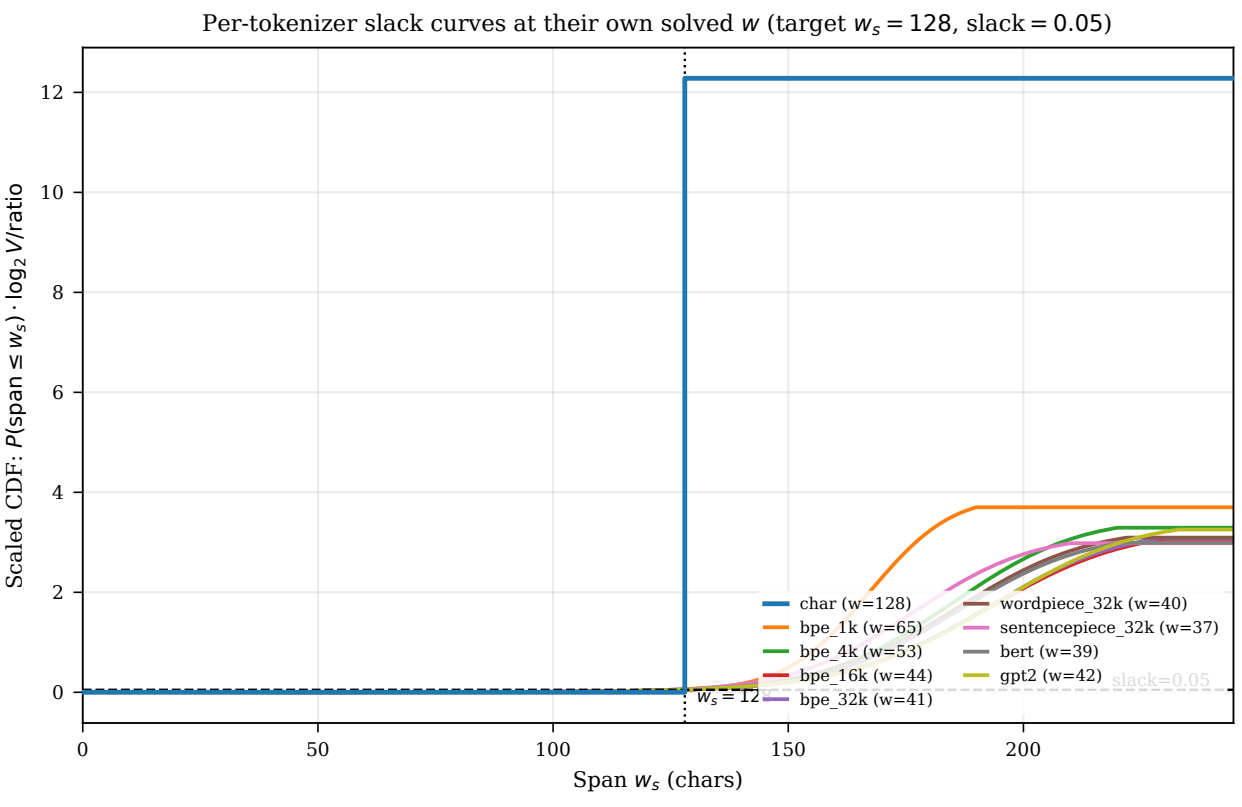
Per-tokenizer slack curves at their own solved w (target $w_s = 128$, slack = 0.05)
Scaled CDF: $P(\text{span} \le w_s) \cdot \log_2 V$/ratio
Span w_s (chars)
$w_s = 128$
slack=0.05
char (w=128)
bpe_1k (w=65)
bpe_4k (w=53)
bpe_16k (w=44)
bpe_32k (w=41)
wordpiece_32k (w=40)
sentencepiece_32k (w=37)
bert (w=39)
gpt2 (w=42)